

Introduction

Spam SMS Detection is a Natural Language Processing (NLP) task that classifies messages into two categories: **spam** or **ham** (not spam). It is widely used in SMS applications and email services to automatically filter unwanted or harmful messages.

In this project, we build a machine learning model to detect spam using the **Multinomial Naive Bayes** algorithm, a popular method for text classification.

Importing Required Libraries

We import essential libraries for data manipulation, visualization, text processing, and machine learning.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import re
```

Load the Dataset

We load the `spam.csv` file and focus only on the two useful columns: `v1` (label) and `v2` (message). We rename them for better clarity.

```
In [6]: df = pd.read_csv("spam.csv", encoding='latin-1')
df = df[['v1', 'v2']]
df.columns = ['label', 'message']
df.head()
```

```
Out[6]:
```

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

Data Cleaning

We clean the dataset by:

- Removing missing values
- Converting labels into numeric format (ham = 0, spam = 1)
- Preprocessing the message text (lowercasing, removing special characters)

```
In [5]: import pandas as pd
import re

# Load the dataset (adjust path if needed)
df = pd.read_csv("spam.csv", encoding='latin-1')

# Keep only necessary columns and rename them
df = df[['v1', 'v2']]
df.columns = ['label', 'message']

# Drop rows with missing values (if any)
df.dropna(inplace=True)

# Convert 'label' column from text to numeric values
df['label_num'] = df['label'].map({'ham': 0, 'spam': 1})

# Define a text preprocessing function
def preprocess_text(text):
    text = text.lower() # Lowercase
    text = re.sub(r'^a-zA-Z0-9', ' ', text) # remove special characters
    text = re.sub(r'\s+', ' ', text) # remove extra spaces
    return text.strip()

# Apply the cleaning function to the 'message' column
df['clean_message'] = df['message'].apply(preprocess_text)

# Display the cleaned data
df[['message', 'clean_message']].head()
```


```
Out[5]:
```

	message	clean_message
0	Go until jurong point, crazy.. Available only ...	go until jurong point crazy available only in ...
1	Ok lar... Joking wif u oni...	ok lar joking wif u oni
2	Free entry in 2 a wkly comp to win FA Cup fina...	free entry in 2 a wkly comp to win fa cup fina...
3	U dun say so early hor... U c already then say...	u dun say so early hor u c already then say
4	Nah I don't think he goes to usf, he lives aro...	nah i don t think he goes to usf he lives arou...

Exploratory Data Analysis

Let's visualize the distribution of spam and ham messages to understand the class balance in the dataset.

```
In [9]: sns.countplot(x='label', data=df)
plt.title("Distribution of Ham and Spam messages")
plt.show()
```



Feature Extraction

We use `CountVectorizer` to convert text data into numerical format that the model can understand. Each SMS becomes a vector of word counts.

```
In [6]: from sklearn.feature_extraction.text import CountVectorizer

# Initialize the CountVectorizer
vectorizer = CountVectorizer()

# Fit the vectorizer on the cleaned messages and transform them into feature vectors
X = vectorizer.fit_transform(df['clean_message'])

# Target variable (Labels: 0 for ham, 1 for spam)
y = df['label_num']

# Display the shape of the resulting feature matrix
print("Feature matrix shape:", X.shape)
```

Feature matrix shape: (5572, 8622)

Train-Test Split

We split the dataset into training and testing sets using an 80-20 ratio to evaluate the model's performance on unseen data.

```
In [11]: # Step 1: Import necessary Libraries
import pandas as pd
import re
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split

# Step 2: Load the dataset (make sure to replace with your actual path)
df = pd.read_csv("spam.csv", encoding='latin-1')

# Step 3: Drop unused or unnamed columns if any
df = df[['v1', 'v2']]
df.columns = ['label', 'message'] # Rename for clarity

# Step 4: Drop any missing values
df.dropna(inplace=True)

# Step 5: Convert labels to numeric (ham = 0, spam = 1)
df['label_num'] = df['label'].map({'ham': 0, 'spam': 1})

# Step 6: Text cleaning function
def preprocess_text(text):
    text = text.lower() # Lowercase
    text = re.sub(r'^a-zA-Z0-9', ' ', text) # remove special characters
    text = re.sub(r'\s+', ' ', text) # remove extra spaces
    return text.strip()

# Step 7: Apply text cleaning
df['clean_message'] = df['message'].apply(preprocess_text)

# Step 8: Feature extraction using CountVectorizer
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(df['clean_message'])

# Step 9: Labels
y = df['label_num']

# Step 10: Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)

# Optional: Check the shape
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

X_train shape: (4457, 8622)
X_test shape: (1115, 8622)
y_train shape: (4457,)
y_test shape: (1115,)
```

Model Training

We use the **Multinomial Naive Bayes** classifier, a common algorithm for classifying text using word frequencies.

```
In [13]: from sklearn.naive_bayes import MultinomialNB
model = MultinomialNB()
model.fit(X_train, y_train)
```

```
Out[13]: ▼ MultinomialNB
MultinomialNB()
```

Making Predictions

Now we use our trained model to predict the labels for the test data.

```
In [17]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Make predictions on the test set
y_pred = model.predict(X_test)

# Print accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Print detailed classification report
print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=['Ham', 'Spam']))
```

Accuracy: 0.97847533632287

Confusion Matrix:

```
[[952  13]
 [ 11 139]]
```

Classification Report:

	precision	recall	f1-score	support
Ham	0.99	0.99	0.99	965
Spam	0.91	0.93	0.92	150
accuracy			0.98	1115
macro avg	0.95	0.96	0.95	1115
weighted avg	0.98	0.98	0.98	1115

Model Evaluation

We evaluate the model using:

- Accuracy Score
- Classification Report (Precision, Recall, F1-score)
- Confusion Matrix

```
In [19]: # Import necessary evaluation modules
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Predict the labels for the test set
y_pred = model.predict(X_test)

# Accuracy score
accuracy = accuracy_score(y_test, y_pred)
print("✅ Accuracy of the Model:", round(accuracy * 100, 2), "%")

# Confusion Matrix
print("\n📊 Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Classification Report
print("\n📄 Classification Report:")
print(classification_report(y_test, y_pred, target_names=['Ham', 'Spam']))
```

✅ Accuracy of the Model: 97.85 %

📊 Confusion Matrix:

```
[[952  13]
 [ 11 139]]
```

📄 Classification Report:

	precision	recall	f1-score	support
Ham	0.99	0.99	0.99	965
Spam	0.91	0.93	0.92	150
accuracy			0.98	1115
macro avg	0.95	0.96	0.95	1115
weighted avg	0.98	0.98	0.98	1115



Conclusion

We successfully built a Spam SMS Detection model using the Multinomial Naive Bayes algorithm. The model performed well and accurately classified SMS messages into spam and ham.



Key Learnings:

- Text data must be cleaned and preprocessed before modeling.
- CountVectorizer is useful for converting text into numbers.
- Naive Bayes works efficiently for spam detection tasks.



Future Improvements:

- Use `TfidfVectorizer` for better feature scaling.
- Remove stopwords for cleaner inputs.
- Try deep learning models for improved accuracy.